

Machine Learning


Reinforcement learning

S. Herbin

`stephane.herbin@onera.fr`


Introduction

Supervised Learning



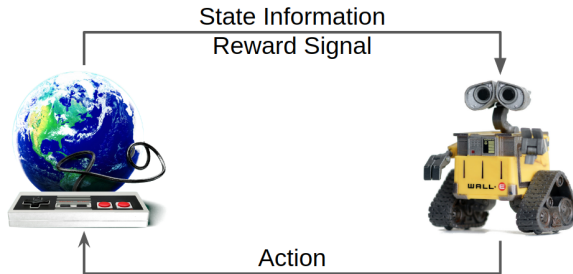
- ▶ Data: Features + Labels
- ▶ Task: Discriminate classes based on features
- ▶ Learning objective: find the good decision that minimizes the expected risk
- ▶ Problem: the expected risk is only available through samples (Generalization issue)

Unsupervised Learning



- ▶ Data: Samples
- ▶ Task: Find structure or good latent representation in samples
- ▶ Learning objective: gather or embed data that share common objective (similarity, pretext task)
- ▶ Problem: what is a good objective?

Reinforcement Learning



- ▶ Data: samples + rewards
 - ▶ Task: Learn how to behave s.t. reward is maximized
 - ▶ Learning objective: generate **sequences** of good decisions
 - ▶ Problem: value of decisions is only known by acting
- ~> need to generate and analyze a large space of possible sequences (exploration vs. exploitation dilemma).

Vocabulary of RL

Action space

Observation space

State space

Reward & return

Policy

Episode, Trajectory

Horizon

Value function

Off policy / In policy

...

RL basics

1. Modeling sequential dependencies, actions and observations
 - ▶ The Markov family
2. Optimizing actions: model based
 - ▶ Dynamic programming
 - ▶ Policy iteration vs. Value iteration
3. Optimizing actions: Reinforcement Learning
 - ▶ Exploration/Exploitation
 - ▶ Fundamental algorithms: SARSA and Q-learning
4. What about Deep learning?
 - ▶ Function approximation
 - ▶ Deep Q-learning: DQN
 - ▶ PPO: Proximal Policy Optimization

Modeling sequential dependencies, actions and observations

Markov Process or Markov Chain

- ▶ Information state: sufficient statistic of history
- ▶ State s_t has Markov property if and only if:

$$p(s_{t+1} \mid s_t) = p(s_{t+1} \mid h_t)$$

where $h_t = (s_0, s_1, \dots, s_t)$ is the history and p is a probability.

↪ Memory-less random process (/walk)

- ▶ Definition of Markov Process $M = (S, P)$
 - ▶ S is a (finite) set of states ($s \in S$)
 - ▶ P is a (stationary) transition model that specifies $p(s_{t+1} = s' \mid s_t = s)$
- ▶ If finite number (N) of states, P is as a matrix $P_{i,j} = p(s_{t+1} = j \mid s_t = i)$. Many theoretical results!!

Markov Reward Process (MRP)

- ▶ Extend Markov Process by rewards
- ▶ Definition of Markov Reward Process (MRP) $M = (S, P, R, \gamma)$
 - ▶ S is a (finite) set of states ($s \in S$)
 - ▶ P is dynamics/transition model that specifies $P(s_{t+1} = s' \mid s_t = s)$
 - ▶ R is a reward function $R(s_t = s)$
- ▶ Note: still no actions
- ▶ If finite number (N) of states, we can express R as a vector

Return & Value Function

- ▶ Episode:
 - ▶ Sequence of states until termination
 - ▶ A “sample”
- ▶ Horizon:
 - ▶ Number of time steps in each episode
 - ▶ Can be finite or infinite

Return & Value Function

- ▶ Episode:
 - ▶ Sequence of states until termination
 - ▶ A “sample”
- ▶ Horizon:
 - ▶ Number of time steps in each episode
 - ▶ Can be finite or infinite
- ▶ Return: G_t (for a MRP)
 - ▶ Discounted sum of rewards from time step t to horizon with factor $\gamma \in [0, 1]$

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

Return & Value Function

- ▶ Episode:
 - ▶ Sequence of states until termination
 - ▶ A “sample”
- ▶ Horizon:
 - ▶ Number of time steps in each episode
 - ▶ Can be finite or infinite
- ▶ Return: G_t (for a MRP)
 - ▶ Discounted sum of rewards from time step t to horizon with factor $\gamma \in [0, 1]$

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

- ▶ State Value Function: $V(s)$ (for a MRP)
 - ▶ Expected return from starting in state s

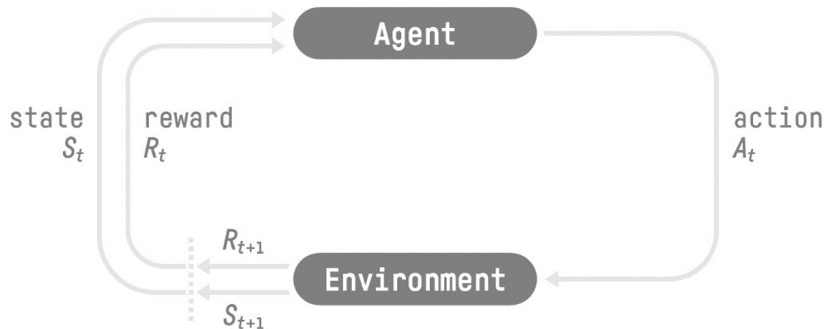
$$V(s) = \mathbb{E}[G_t \mid s_t = s] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots \mid s_t = s]$$

Discount Factor

$$V(s) = \mathbb{E}[G_t \mid s_t = s] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots \mid s_t = s]$$

- ▶ Mathematically convenient (avoid infinite returns and values)
- ▶ Humans often act as if there's a discount factor $\gamma < 1$ on impact to future
- ▶ $\gamma = 0$: Only care about immediate reward
- ▶ $\gamma = 1$: Future reward is as beneficial as immediate reward
- ▶ If episode lengths are always finite, can use $\gamma = 1$ (but don't have to)

The Reinforcement loop



Markov Decision Process (MDP)

- ▶ Markov Decision Process is Markov Reward Process + actions
- ▶ Definition of MDP
 - ▶ S is a (finite) set of Markov states $s \in S$
 - ▶ A is a (finite) set of actions $a \in A$
 - ▶ P is dynamics/transition model for each action, that specifies $P(s_{t+1} = s' \mid s_t = s, a_t = a)$
 - ▶ R is a reward function $R(s_t = s, a_t = a)$ possibly random
 - ▶ Sometimes R is also defined based on (s) or on (s, a, s')
 - ▶ Discount factor $\gamma \in [0, 1]$
- ▶ MDP is tuple (S, A, P, R, γ)

MDP (cont'd)

- ▶ MDP is tuple (S, A, P, R, γ)
- ▶ Optional components:
 - ▶ $\rho_0 : S \rightarrow \mathbb{R}^+$: a distribution of start states
 - ▶ uniform distribution: the agent can start in any state – implicit assumption of MDP definition above
 - ▶ non-uniform distribution: the agent starts its episodes in only some of the states; e.g., it's unlikely that a game will start in a terminal state

MDP (cont'd)

- ▶ MDP is tuple (S, A, P, R, γ)
- ▶ Optional components:
 - ▶ $\rho_0 : S \rightarrow \mathbb{R}^+$: a distribution of start states
 - ▶ uniform distribution: the agent can start in any state – implicit assumption of MDP definition above
 - ▶ non-uniform distribution: the agent starts its episodes in only some of the states; e.g., it's unlikely that a game will start in a terminal state
 - ▶ $T \subset S$: set of terminal states
 - ▶ important for episodic MDPs
 - ▶ or if there is not fixed horizon, but the episodes should be finite

MDP (cont'd)

- ▶ MDP is tuple (S, A, P, R, γ)
- ▶ Optional components:
 - ▶ $\rho_0 : S \rightarrow \mathbb{R}^+$: a distribution of start states
 - ▶ uniform distribution: the agent can start in any state – implicit assumption of MDP definition above
 - ▶ non-uniform distribution: the agent starts its episodes in only some of the states; e.g., it's unlikely that a game will start in a terminal state
 - ▶ $T \subset S$: set of terminal states
 - ▶ important for episodic MDPs
 - ▶ or if there is not fixed horizon, but the episodes should be finite
 - ▶ γ : discount factor
 - ▶ important to quantify the importance of future
 - ▶ some treat γ as a hyperparameter and not part of the definition
 - ~> different optimal policies can be found
 - ~> depends on how the optimal policy is defined

MDP Policy

- ▶ Policy specifies what action to take in each state
 - ▶ Can be deterministic or stochastic
- ▶ For generality, considered as a stationary conditional distribution

$$\pi(a \mid s) = P(a_t = a | s_t = s)$$

- ▶ Other name for a control law.
- ▶ This is what is expected to be learned!

MDP + Policy = MRP

- ▶ MDP + Policy $\pi(a | s)$ = Markov Reward Process
- ▶ $[s_t, a_{t-1}]$ is Markov
- ▶ Precisely, it is the MRP $(S, R^\pi, P^\pi, \gamma)$ where

$$R^\pi(s) = \sum_{a \in A} \pi(a | s) R(s, a)$$

$$P^\pi(s' | s) = \sum_{a \in A} \pi(a | s) P(s' | s, a)$$

- ▶ Implies we can use same techniques to evaluate the value of a policy for an MDP as we could to compute the value of a MRP, by defining a MRP with R^π and P^π

MDP and Value functions

- ▶ Definition of Return G_t (same as MRP)
 - ▶ Discounted sum of rewards from time step t to horizon

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

MDP and Value functions

- ▶ Definition of Return G_t (same as MRP)

- ▶ Discounted sum of rewards from time step t to horizon

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

- ▶ Definition of *On-Policy* State Value Function $V^\pi(s)$: depends on policy π !

- ▶ Expected return from starting in state s under policy π

$$V^\pi(s) = \mathbb{E}_\pi[G_t \mid s_t = s] \tag{1}$$

$$= \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots \mid s_t = s] \tag{2}$$

MDP and Value functions

- ▶ Definition of Return G_t (same as MRP)

- ▶ Discounted sum of rewards from time step t to horizon

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

- ▶ Definition of *On-Policy* State Value Function $V^\pi(s)$: depends on policy π !

- ▶ Expected return from starting in state s under policy π

$$V^\pi(s) = \mathbb{E}_\pi[G_t \mid s_t = s] \tag{1}$$

$$= \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots \mid s_t = s] \tag{2}$$

- ▶ Definition of *On-Policy* State-Action Value Function $Q^\pi(s, a)$

- ▶ Expected return from starting in state s , taking action a and then following policy π

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t \mid s_t = s, a_t = a]$$

$$= \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots \mid s_t = s, a_t = a]$$

Other Markov models

- ▶ In MRP and MDP, the state (i.e. the environment) is fully **observable**. In many problems, this is not possible.
- ▶ HMM: Hidden Markov Model
 - ▶ The state s is partially observable from o with conditional probability: $p(o | s)$
- ▶ POMDP: Partially Observable Markov Decision Process
 - ▶ MDP + HMM

Optimizing actions: model based approach

Optimal value functions

- ▶ The Optimal Value Function, $V^*(s)$ gives the expected return if you start in state s and always act according to the optimal policy in the environment:

$$V^*(s) = \max_{\pi} \mathbb{E}_{\pi}[G_t \mid s_t = s]$$

- ▶ The Optimal Action-Value Function, $Q^*(s, a)$ gives the expected return if you start in state s , take an arbitrary action a , and then forever after act according to the optimal policy in the environment:

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{\pi}[G_t \mid s_t = s, a_t = a]$$

- ▶ The goal is to compute the optimal policy π^* that maximizes the expected return:

$$\pi^*(s) \in \arg \max_{\pi} V^{\pi}(s)$$

MDP optimality

It can be shown that:

- ▶ There exists a unique optimal value function $V^*(s)$
- ▶ In an infinite horizon problem (i.e. agents act forever), there exists an optimal policy π^* for an MDP, that is
 - ▶ deterministic
 - ▶ stationary (does not depend on time step)
 - ▶ unique? \rightsquigarrow Not necessarily, may have state-actions with identical optimal values.
 - ▶ The optimal policy can be expressed as:

$$\pi^*(s) \in \arg \max_a Q^*(s, a)$$

- ▶ In an finite horizon problem undiscounted return ($\gamma = 1$) needs to make time as an argument \rightsquigarrow trade-off between time and impact of actions

Bellman Equation and Bellman Backup Operators

The fundamental equations!

- ▶ Value functions of a policy must satisfy the Bellman equations

$$V^\pi(s) = \sum_a \pi(a | s) \left[R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) V^\pi(s') \right]$$
$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) \sum_{a'} \pi(a' | s') Q^\pi(s', a')$$

- ▶ *“The value of your starting point is the reward you expect to get from being there, plus the value of wherever you land next.”*
- ▶ Optimal equations:

$$V^*(s) = \max_a \left[R(s, a) + \gamma \sum_{s' \in S} P^\pi(s' | s, a) V^*(s') \right]$$
$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) \max_{a'} Q^*(s', a')$$

Bellman Backup Operators

Bellman backup operator B^π

$$B^\pi V(s) = \sum_a \pi(a | s) \left[R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) V(s') \right]$$

Bellman optimal backup operator B^*

$$B^* V(s) = \max_a \left[R(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V(s') \right]$$

- ▶ Read B^* and B^π as operators applied to V and yielding a value function over all states s
- ▶ Need to know transition probability $p(s' | s, a)$ and rewards function $R(s, a)$ to compute it.

MDP: Computing Optimal Policy and Optimal Value I

Two strategies

1. Learn the Policy by iteration

- ▶ Start with an approximate or random policy
- ▶ Evaluate it
- ▶ Improve it

2. Learn the Value function by iteration

- ▶ Idea: Maintain optimal value of starting in a state s if we have a finite number of steps k left in the episode
- ▶ Iterate to consider longer and longer episodes
- ▶ Define the policy from the Value function V

3. Estimate the value by random sampling (Monte-Carlo)

- ▶ Idea (simple): generate episodes, collect rewards and returns, compute average return.

MDP Policy Evaluation, Iterative Algorithm

- ▶ Goal: For a given π , determine V^π
- ▶ iterative approach:
 - ▶ Initialize $V_0(s) = 0$ for all s
 - ▶ For $k = 1$ until convergence
 - ▶ For all s in S :

$$V_k^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in S} p(s' | s, \pi(s)) V_{k-1}^\pi(s')$$

- ▶ This is a Bellman backup for a particular policy

MDP Policy Evaluation

Algorithm 1: Policy Evaluation

Input: MDP, policy π , small positive number θ

Output: $V \approx v_\pi$

Initialize V arbitrarily (e.g., $V(s) = 0$ for all $s \in \mathcal{S}^+$)

repeat

$\Delta \leftarrow 0$

for $s \in \mathcal{S}$ **do**

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a | s) \sum_{s'} p(s' | s, a)(r + \gamma V(s'))$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

end

until $\Delta < \theta$;

return V

Policy Evaluation as Bellman Operations

- ▶ Bellman backup operator B^π for a particular policy is defined as

$$B^\pi V(s) = R^\pi(s) + \gamma \sum_{s' \in S} P^\pi(s' | s) V(s')$$

- ▶ Policy evaluation amounts to computing the fixed point of B^π
- ▶ To do policy evaluation, repeatedly apply operator until V stops changing

$$V^\pi = B^\pi B^\pi B^\pi B^\pi \dots B^\pi V$$

- ▶ It converges because B^π is contractive for $\gamma < 1$.

Policy Improvement

- ▶ Act greedily with respect to V^π to choose the action.
- ▶ Compute state-action value of a policy π_k
 - ▶ For s in S and a in A :

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) V^\pi(s')$$

- ▶ Compute new policy π_{k+1} for all $s \in S$

$$\pi_{k+1}(s) \in \arg \max_{a \in A} Q^{\pi_k}(s, a)$$

Policy Improvement



Algorithm 2: Policy Improvement

Input: MDP, value function V

Output: policy π'

```
for  $s \in S$  do
  for  $a \in A$  do
     $Q(s, a) \leftarrow \sum_{s'} p(s' | s, a)(r + \gamma V(s'))$ 
  end
   $\pi'(s) \leftarrow \arg \max_{a \in A(s)} Q(s, a)$ 
end
return  $\pi'$ 
```

MDP Policy Iteration



MDP Policy Iteration

Algorithm 3: Policy Iteration

Input: MDP, small positive number θ

Output: policy $\pi \approx \pi^*$

Initialize π arbitrarily (e.g., $\pi(a | s) = \frac{1}{|A|}$ for all $s \in S$ and $a \in A$)

policy-stable \leftarrow *false*

repeat

$V \leftarrow \text{Policy_Evaluation}(\text{MDP}, \pi, \theta)$

$\pi' \leftarrow \text{Policy_Improvement}(\text{MDP}, V)$

if $\pi = \pi'$ **then**

 | *policy-stable* \leftarrow *true*

end

$\pi \leftarrow \pi'$

until *policy-stable* = *true*;

return π

Policy iteration stopping condition

- ▶ Since $\pi_{k+1}(s) \in \arg \max_{a \in A} Q^{\pi_k}(s, a)$ we can ensure that

$$\forall s \in S, V^{\pi_{k+1}}(s) \geq V^{\pi_k}(s)$$

- ▶ Iterations stop for a k where

$$Q^{\pi}(s, \pi_{k+1}(s)) = \max_a Q^{\pi}(s, a) = Q^{\pi}(s, \pi_k(s)) = V^{\pi_k}(s)$$

- ▶ The iterations stop when Bellman optimality condition is verified

$$\forall s \in S, V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

- ▶ The output policy π_k is optimal ($\pi_k = \pi^*$)

Value Iteration (VI)

- ▶ Set $k = 1$
- ▶ Initialize $V_0(s) = 0$ for all states s
- ▶ Loop until convergence
 - ▶ For each state s

$$V_{k+1}(s) = \max_{a \in A} R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) V_k(s')$$

- ▶ Interpreted as Bellman backup on value function (fixed point iteration)

$$V_{k+1} = B^\pi V_k$$

- ▶ The final policy is defined as:

$$\pi(s) \in \arg \max_{a \in A} \left[R(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) V(s') \right]$$

Algorithm 4: Value Iteration

Input: MDP, small positive number θ

Output: policy $\pi \approx \pi_*$

Initialize V arbitrarily (e.g., $V(s) = 0$ for all $s \in \mathcal{S}^+$)

repeat

$\Delta \leftarrow 0$

for $s \in \mathcal{S}$ **do**

$v \leftarrow V(s)$

$V(s) \leftarrow \max_{a \in A} \sum_{s'} p(s' | s, a)(r + \gamma V(s'))$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

end

until $\Delta < \theta$;

$\pi \leftarrow \text{Policy_Improvement}(\text{MDP}, V)$

return π

Evaluate the Policy by random sampling + iteration (Monte-Carlo)

- ▶ Generate a series of episodes $[s_0, a_1, s_1, r_1, \dots, s_T, a_T, r_T]$ using current policy π_k , and estimate the partial return $G_t = \sum_{j=t}^T \gamma^{j-t} r_j$
- ▶ Value = mean return on each state (Law of Large Numbers) for first-visit or every-visit state
- ▶ Improve policy

Monte-Carlo policy evaluation

Algorithm 5: First-Visit MC Prediction (*for action values*)

Input: policy π , positive integer num_episodes

Output: value function Q ($\approx q_\pi$ if num_episodes is large enough)

Initialize $N(s, a) = 0$ for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Initialize returns_sum(s, a) = 0 for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

for $i \leftarrow 1$ **to** num_episodes **do**

 Generate an episode $S_0, A_0, R_1, \dots, S_T$ using π

for $t \leftarrow 0$ **to** $T - 1$ **do**

if (S_t, A_t) is a first visit (with return G_t) **then**

$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$

 returns_sum(S_t, A_t) \leftarrow returns_sum(S_t, A_t) + G_t


end

end

$Q(s, a) \leftarrow$ returns_sum(s, a) / $N(s, a)$ for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

return Q

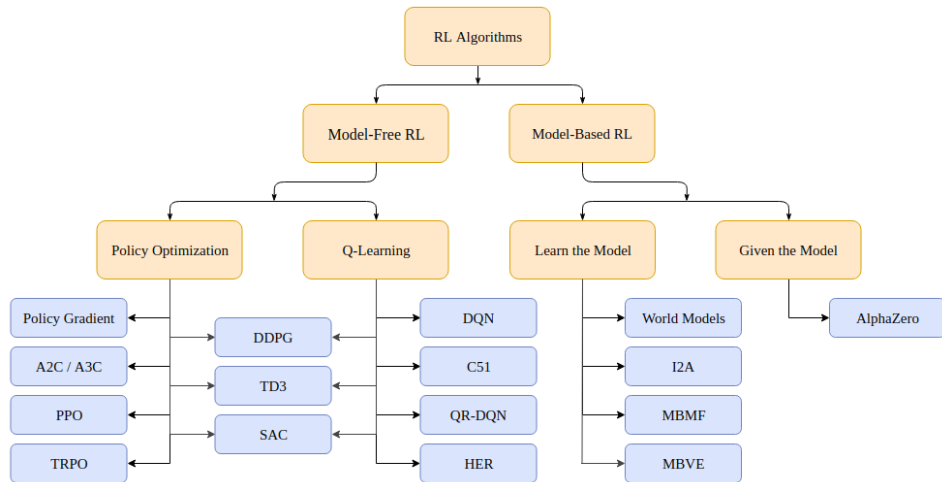
Small demo



Grid World

Learning policy

The huge family of RL algorithms



Learning without models

- ▶ Assumption: We don't have the exact model of the environment (i.e., model-free), but we can query the environment ("playing roll-outs")
 - ▶ state space and action space are in principle known
 - ▶ we don't know the transition probabilities beforehand
 - ▶ we don't know the reward distribution beforehand

Learning without models

- ▶ Assumption: We don't have the exact model of the environment (i.e., model-free), but we can query the environment ("playing roll-outs")
 - ▶ state space and action space are in principle known
 - ▶ we don't know the transition probabilities beforehand
 - ▶ we don't know the reward distribution beforehand
- ▶ Remarks:
 - ▶ If we would know the MDP, we only have to do "planning" to find the optimal policy
 - ▶ If we first learn the MDP and then apply planning to the learned MDP, we do "model-based" RL

Learning without models

- ▶ Assumption: We don't have the exact model of the environment (i.e., model-free), but we can query the environment ("playing roll-outs")
 - ▶ state space and action space are in principle known
 - ▶ we don't know the transition probabilities beforehand
 - ▶ we don't know the reward distribution beforehand
- ▶ Remarks:
 - ▶ If we would know the MDP, we only have to do "planning" to find the optimal policy
 - ▶ If we first learn the MDP and then apply planning to the learned MDP, we do "model-based" RL
- ▶ Goal: We want to learn $V^\pi(s)$ or $Q^\pi(s, a)$ (depending on the RL algorithm we want to use) by only querying the unknown MDP

Exploration & exploitation

When the model + environment are unknown, we need to find ways to understand how they behave.

Reinforcement learning is like trial-and-error learning

- ▶ The agent should discover a good policy
- ▶ From its experiences of the environment
- ▶ Without losing too much reward along the way

Two design principles:

- ▶ exploration = generate relevant experience of unknown environment
- ▶ exploitation = use experience to efficiently improve the policy π (= maximize the reward)

RL first basic ideas

Greedy Policy iteration may be inefficient:

- ▶ If π is deterministic, we may not observe all possible actions $a \in A$ in a state s
- ▶ So, we cannot compute $Q(s, a)$ for any $a \neq \pi(s)$

↪ How to interleave policy evaluation and improvement?

Two tools

- ▶ Randomness in decision: ϵ -greediness
- ▶ Sequential stochastic approximation

ϵ -greedy Policies

Simplest idea for ensuring continual exploration

- ▶ All m actions are tried with non-zero probability
- ▶ With probability $1 - \epsilon$ choose the greedy action
- ▶ With probability ϵ choose an action at random

$$\pi(a \mid s) = \begin{cases} 1 - \epsilon + \epsilon/m & \text{if } a = \arg \max_{a' \in A} Q(s, a') \\ \epsilon/m & \text{otherwise} \end{cases}$$

One can show that

- ▶ For any ϵ -greedy policy π' , the ϵ -greedy policy wrt Q^π is a monotonic improvement of the value function

$$V^{\pi'} \geq V^\pi$$

Greedy in the Limit of Infinite Exploration (GLIE)

Definition of GLIE:

- ▶ All state-action pairs are visited an infinite number of times
- ▶ Behavior policy (policy used to act in the world) converges to greedy policy

$$\lim_{i \rightarrow \infty} \pi(a \mid s) \rightarrow \arg \max_{a \in A} Q(s, a)$$

with probability 1

Simple Strategy:

- ▶ ϵ -greedy where ϵ is annealed close to 0 with $\epsilon_i = 1/i$

Theorem:

- ▶ GLIE Monte-Carlo control converges to the optimal state-action value function
 $Q(s, a) \rightarrow Q^*(s, a)$

Monte-Carlo policy optimization

Algorithm 6: First-Visit GLIE MC Control

Input: positive integer num_episodes, GLIE $\{\epsilon_i\}$

Output: policy π ($\approx \pi_*$ if num_episodes is large enough)

Initialize $Q(s, a) = 0$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Initialize $N(s, a) = 0$ for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

for $i \leftarrow 1$ **to** num_episodes **do**

$\epsilon \leftarrow \epsilon_i$

$\pi \leftarrow \epsilon$ -greedy(Q)

 Generate an episode $S_0, A_0, R_1, \dots, S_T$ using π

for $t \leftarrow 0$ **to** $T - 1$ **do**

if (S_t, A_t) is a first visit (with return G_t) **then**

$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)}(G_t - Q(S_t, A_t))$

end

end

return π

Model-free Policy Iteration with TD Methods

- ▶ Monte-Carlo approach requires to sample full episodes before updating the policy using the return G_t
- ▶ One idea: use the current estimate of the return $r_{t+1} + \gamma Q(s_{t+1}, a_{t+1})$ instead of G_t and update at each time step
= this is *Temporal Difference* learning
- ▶ It allows the exploitation of incomplete sequences

This gives the *SARSA* algorithm.

Algorithm 7: Sarsa

Input: policy π , positive integer num_episodes, small positive fraction α , GLIE $\{\epsilon_i\}$

Output: value function Q ($\approx q_\pi$ if num_episodes is large enough)

Initialize Q arbitrarily (e.g., $Q(s, a) = 0$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$, and $Q(\text{terminal-state}, \cdot) = 0$)

for $i \leftarrow 1$ **to** num_episodes **do**

$\epsilon \leftarrow \epsilon_i$

 Observe S_0

 Choose action A_0 using policy derived from Q (e.g., ϵ -greedy)

$t \leftarrow 0$

repeat

 Take action A_t and observe R_{t+1}, S_{t+1}

 Choose action A_{t+1} using policy derived from Q (e.g., ϵ -greedy)

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$

$t \leftarrow t + 1$

until S_t is terminal;

end

return Q

Convergence Properties of SARSA

- ▶ Theorem:
SARSA for finite-state and finite-action MDPs converges to the optimal action-value, $Q(s, a) \rightarrow Q^*(s, a)$, under the following conditions:
 1. The policy sequence $\pi_t(a | s)$ satisfies the condition of GLIE
 2. The step-sizes α_t satisfy the **Robbins-Munro sequence** such that

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

- ▶ For example $\alpha_t = \frac{1}{t}$ satisfies the above condition
- ▶ In practice, α is often kept constant.

Q-Learning: Learning the Optimal State-Action Value

- ▶ A possible alternative is to directly estimate the value of π^* while acting with another behavior policy π_b ?
- ▶ This is called an *off-policy* RL algorithm
- ▶ Key idea: Maintain state-action Q estimates and use it to bootstrap: use the value of the best future action
- ▶ This gives the *Q-Learning* updating step:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha((r_{t+1} + \gamma \max_{a' \in A} Q(s_{t+1}, a')) - Q(s_t, a_t))$$

- ▶ Same convergence property to optimal policy as SARSA (GLIE + Robbins-Munro)

Algorithm 8: Q-Learning (Sarsamax)

Input: policy π , positive integer num_episodes, small positive fraction α , GLIE $\{\epsilon_i\}$

Output: value function Q ($\approx q_\pi$ if num_episodes is large enough)

Initialize Q arbitrarily (e.g., $Q(s, a) = 0$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$, and $Q(\text{terminal-state}, \cdot) = 0$)

for $i \leftarrow 1$ **to** num_episodes **do**

$\epsilon \leftarrow \epsilon_i$

 Observe S_0

$t \leftarrow 0$

repeat

 Choose action A_t using policy derived from Q (e.g., ϵ -greedy)

 Take action A_t and observe R_{t+1}, S_{t+1}

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$

$t \leftarrow t + 1$

until S_t is terminal;

end

return Q

Policies and experience


- ▶ On-policy learning
 - ▶ “Learn on the job”
 - ▶ Learn about policy π from experience sampled from π
 - ▶ SARSA is one example
- ▶ Off-policy learning
 - ▶ “Look over someone’s shoulder”
 - ▶ Learn about policy π from experience sampled from another one (μ)
 - ▶ Q-learning is one example

RL summary

<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Iterative Policy Evaluation	TD Learning
$V(s) \leftarrow \mathbb{E}[R + \gamma V(S') \mid s]$	$V(S) \stackrel{\alpha}{\leftarrow} R + \gamma V(S')$
Q-Policy Iteration	Sarsa
$Q(s, a) \leftarrow \mathbb{E}[R + \gamma Q(S', A') \mid s, a]$	$Q(S, A) \stackrel{\alpha}{\leftarrow} R + \gamma Q(S', A')$
Q-Value Iteration	Q-Learning
$Q(s, a) \leftarrow \mathbb{E}\left[R + \gamma \max_{a' \in \mathcal{A}} Q(S', a') \mid s, a\right]$	$Q(S, A) \stackrel{\alpha}{\leftarrow} R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')$

(from D. Silver slides)

Small demo - RL



Grid World

Deep RL

Deep Learning in RL

- ▶ High dimensional (continuous) states: e.g. images, kinematics of articulated structures
- ▶ Parametric representation of the policy

$$\hat{Q}(s, a; w) \approx Q(s, a)$$

- ▶ Function Approximations techniques (gradient descent on relevant MSE loss) and generalization issues

Incremental Model-Free Control Approaches

- ▶ Similar to policy evaluation, true state-action value function for a state is unknown and we substitute an estimate
- ▶ In Monte Carlo methods, use a return G_t as a substitute target

$$\Delta w = \alpha(G_t - \hat{Q}(s_t, a_t; w)) \nabla_w \hat{Q}(s_t, a_t; w)$$

- ▶ For SARSA instead use a TD target $r + \gamma \hat{Q}(s', a'; w)$ which leverages the current function approximations value

$$\Delta w = \alpha(r + \gamma \hat{Q}(s', a'; w) - \hat{Q}(s, a; w)) \nabla_w \hat{Q}(s, a; w)$$

- ▶ For Q-learning instead use a TD target $r + \gamma \max_{a'} \hat{Q}(s', a'; w)$ which leverages the max of the current function approximations value

$$\Delta w = \alpha(r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \hat{Q}(s, a; w)) \nabla_w \hat{Q}(s, a; w)$$

Using these Ideas to do Deep RL for the Atari challenge

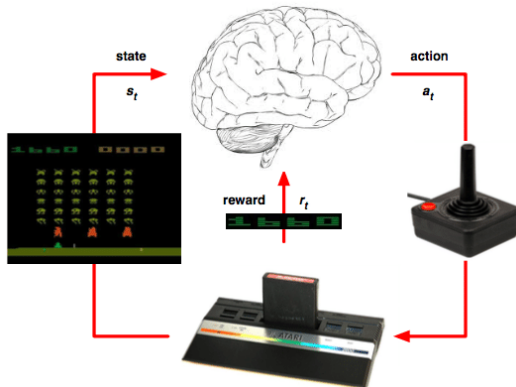



Image by David Silver

Using these Ideas to do Deep RL in Atari

- ▶ End-to-end learning of values $Q(s, a)$ from pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Output is $Q(s, a)$ for 18 joystick/button positions
- ▶ Reward is change in score for that step
- ▶ Network architecture and hyperparameters fixed across all games



DQN source code: <https://github.com/deepmind/dqn>

Q-Learning with Value Function Approximation

- ▶ Minimize MSE loss by stochastic gradient descent
- ▶ Converges to the optimal $Q^*(s, a)$ using **table lookup** representation
- ▶ However Q-learning with Value Function Approximation is unstable
- ▶ Two of the issues causing problems:
 - ▶ Correlations between samples violates i.i.d assumption of DNNs
 - ▶ Non-stationary targets
- ▶ Deep Q-learning (DQN) addresses both of these challenges by
 - ▶ Experience replay
 - ▶ Fixed Q-targets

DQNs: Replay Buffer

- ▶ To help remove correlations, store dataset (called a **replay buffer**) \mathcal{D} from prior experience
- ▶ To perform experience replay, repeat the following:
 1. $(s, a, r, s') \sim \mathcal{D}$: sample experience tuple from the dataset
 2. Compute the target value for the sampled s : $r + \gamma \max_{a'} \hat{Q}(s', a'; w)$
 3. Use stochastic gradient descent to update the network weights

$$\Delta w = \alpha(r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \hat{Q}(s, a; w)) \nabla_w \hat{Q}(s, a; w)$$

- ▶ Remarks:
 - ▶ Fixed sized buffer \rightsquigarrow first-in–first-out scheme (as default implementation)
 - ▶ heuristic trade-off between performing new episodes and sampling from the replay buffer
- \rightsquigarrow Can treat the target as a scalar, but the weights will get updated on the next round, changing the target value

DQNs: Fixed Q-Targets

- ▶ To help improve stability, fix the **target weights** used in the target calculation for multiple updates
- ▶ Target network uses a different set of weights than the weights being updated
- ▶ Let parameters w^- be the set of weights used in the target and w be the weights that are being updated
- ▶ Slight change to computation of target value:
 - ▶ $(s, a, r, s') \sim \mathcal{D}$: sample experience tuple from the dataset
 - ▶ Compute the target value for the sampled s : $r + \gamma \max_{a'} \hat{Q}(s', a'; w^-)$
 - ▶ Use stochastic gradient descent to update the network weights


$$\Delta w = \alpha(r + \gamma \max_{a'} \hat{Q}(s', a'; w^-) - \hat{Q}(s, a; w)) \nabla_w \hat{Q}(s, a; w)$$

- ▶ Remark:
 - ▶ Hyperparameter how often you update w^-
 - ▶ Trade-off between updating too often (\rightsquigarrow instability) and too rarely (\rightsquigarrow too old state information)

DQN Summary

- ▶ DQN uses experience replay and fixed Q-targets
- ▶ Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory \mathcal{D}
- ▶ Sample random mini-batch of transitions (s, a, r, s') from \mathcal{D}
- ▶ Compute Q-learning targets wrt old, fixed parameters w^-
 - ▶ Update w^- from time to time
- ▶ Optimizes MSE between Q-network and Q-learning targets
- ▶ Uses stochastic gradient descent

Actor-Critic algorithms I



Actor-Critic algorithms II

- ▶ Goal: better balance between exploration and exploitation.
- ▶ **Actor (Policy)**: Maintains a policy $\pi(a|s; \theta)$.
- ▶ **Critic (Value Function)**: Learns a value function $V(s; \phi)$. Estimates the expected cumulative reward.
- ▶ **Interaction and Learning**:
 1. The Actor interacts with the environment.
 2. The Critic updates its value function.
 3. The Actor updates its policy.

PPO: a popular actor-critic strategy

- ▶ **Proximal Policy Optimization.**
- ▶ **On-Policy:** Learns the policy by interacting with the environment using that same policy.
- ▶ **Trust Region:** Improves the policy while staying *close* to the previous policy. Essential for stability.
- ▶ **Simple Implementation:** Relatively easy to implement and tune.

PPO Objective Function I

$$L(\theta) = \mathbb{E}_t [\min (r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

- ▶ $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ **(Probability Ratio)**: Measures how much the probability of taking action a_t in state s_t has changed between the current policy π_θ and the old policy $\pi_{\theta_{old}}$. A value close to 1 indicates a small change.
- ▶ A_t **(Advantage Function)**: Quantifies how much better taking action a_t in state s_t was compared to the average return expected in that state. It helps the algorithm learn which actions are truly beneficial. A common way to calculate it is: $A_t = R_t + \gamma V(s_{t+1}) - V(s_t)$, where R_t is the reward received at time t , γ is the discount factor, and $V(s)$ is the value function.

PPO Objective Function II

$$L(\theta) = \mathbb{E}_t [\min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

- ▶ **ϵ (Clipping Parameter):** A hyperparameter that defines the size of the trust region. It limits how much the policy can change in a single update. Typical values are around 0.1 or 0.2. It prevents the policy from making overly large updates that could destabilize training.
- ▶ **$\text{clip}(x, 1 - \epsilon, 1 + \epsilon)$ (Clipping Function):** Limits the probability ratio $r_t(\theta)$ to the range $[1 - \epsilon, 1 + \epsilon]$. This is the key mechanism for enforcing the trust region.
- ▶ **$\min(\cdot, \cdot)$ (Minimum Function):** Takes the minimum of the two arguments. This ensures that the policy update is both beneficial (positive advantage) and within the trust region.

How the Critic Learns


- ▶ **Goal:** Accurately estimate the value function $V(s)$ (expected cumulative reward from state s).
- ▶ **Method:** Typically Temporal Difference (TD) learning (e.g., TD(0)) combined with function approximation (often a neural network).
- ▶ **Process (Simplified):**
 1. Collect experiences (s, a, r, s') from the Actor.
 2. Calculate a target value (e.g., using TD(0): $r + \gamma V(s')$).
 3. Update the Critic's parameters (e.g., neural network weights) to minimize the difference between the current value estimate $V(s)$ and the target.
- ▶ **Advantages:**
 - ▶ **Stability:** The clipping mechanism makes PPO more stable.
 - ▶ **Sample Efficiency:** Learns good policies with fewer interactions.
 - ▶ **Ease of Implementation:** Relatively simple to implement.

Other approaches

MANY!

(See references)

Small demo - Deep RL



Puck World

Références

Many of the slides come from these references

- ▶ <https://www.davidsilver.uk/teaching/>
- ▶ <https://web.stanford.edu/class/cs234/modules.html>
- ▶ https://github.com/automl-edu/RL_lecture